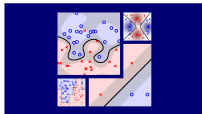


Machine Learning Techniques (機器學習技法)



Lecture 13: Deep Learning

Hsuan-Tien Lin (林軒田)

htlin@csie.ntu.edu.tw

Department of Computer Science
& Information Engineering

National Taiwan University
(國立台灣大學資訊工程系)



Roadmap

- 1 Embedding Numerous Features: Kernel Models
- 2 Combining Predictive Features: Aggregation Models
- 3 Distilling Implicit Features: Extraction Models


Lecture 12: Neural Network

automatic **pattern feature extraction** from **layers of neurons** with **backprop** for GD/SGD

Lecture 13: Deep Learning

- Deep Neural Network
- Autoencoder
- Denoising Autoencoder
- Principal Component Analysis

Physical Interpretation of NNet Revisited



- each layer: **pattern feature extracted** from data, **remember? :-)**
- how many neurons? how many layers?
—more generally, **what structure?**
 - subjectively, **your design!**
 - objectively, **validation, maybe?**

structural decisions:
key issue for applying NNet

Shallow versus Deep Neural Networks

shallow: few (hidden) layers; deep: many layers

Shallow NNet


- more **efficient** to train (○)
- **simpler** structural decisions (○)
- theoretically **powerful enough** (○)

Deep NNet

- **challenging** to train (×)
- **sophisticated** structural decisions (×)
- **'arbitrarily' powerful** (○)
- more **'meaningful'?** (see next slide)

deep NNet (**deep learning**)
gaining attention in recent years

Meaningfulness of Deep Learning



- 'less burden' for each layer: **simple** to **complex** features
- natural for **difficult** learning task with **raw features**, like **vision**

deep NNet: currently popular in
vision/speech/...

Challenges and Key Techniques for Deep Learning


- difficult **structural decisions**:
 - subjective with **domain knowledge**: like **convolutional NNet** for images
- high **model complexity**:
 - no big worries if **big enough data**
 - **regularization** towards noise-tolerant: like
 - **dropout** (tolerant when network corrupted)
 - **denoising** (tolerant when input corrupted)
- hard **optimization problem**:
 - **careful initialization** to avoid bad local minimum: called **pre-training**
- huge **computational complexity** (worsen with **big data**):
 - novel hardware/architecture: like **mini-batch with GPU**

IMHO, careful **regularization** and **initialization** are key techniques

A Two-Step Deep Learning Framework

Simple Deep Learning

- 1 for $\ell = 1, \dots, L$, **pre-train** $\{w_{ij}^{(\ell)}\}$ assuming $w_*^{(1)}, \dots, w_*^{(\ell-1)}$ fixed



- 2 **train with backprop** on **pre-trained** NNet to **fine-tune** all $\{w_{ij}^{(\ell)}\}$

will focus on **simplest pre-training** technique
along with **regularization**

Fun Time

For a deep NNet for written character recognition from raw pixels, which type of features are more likely extracted after the first hidden layer?

- ① pixels
- ② strokes
- ③ parts
- ④ digits

Fun Time

For a deep NNet for written character recognition from raw pixels, which type of features are more likely extracted after the first hidden layer?

- ① pixels
- ② strokes
- ③ parts
- ④ digits



Reference Answer: ②

Simple strokes are likely the 'next-level' features that can be extracted from raw pixels.

Information-Preserving Encoding

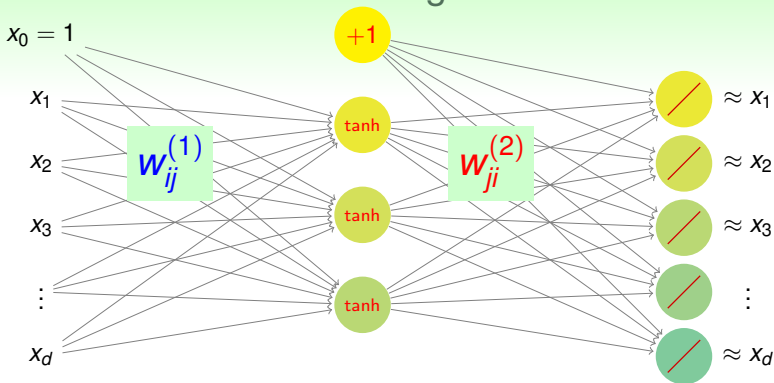
- **weights**: feature transform, i.e. **encoding**
- **good weights**: information-preserving encoding
—next layer same info. with different representation
- **information-preserving**:

decode accurately after encoding



idea: **pre-train weights** towards
information-preserving encoding

Information-Preserving Neural Network



- autoencoder:** $d \rightarrow \tilde{d} \rightarrow d$ NNet with goal $g_i(\mathbf{x}) \approx x_i$
—learning to **approximate identity function**
- $w_{ij}^{(1)}$: encoding weights; $w_{ji}^{(2)}$: decoding weights

why **approximating identity function**?

Usefulness of Approximating Identity Function

if $\mathbf{g}(\mathbf{x}) \approx \mathbf{x}$ using some **hidden** structures on the **observed data** \mathbf{x}_n

- for supervised learning:
 - **hidden structure** (essence) of \mathbf{x} can be used as **reasonable transform** $\Phi(\mathbf{x})$—learning **‘informative’ representation** of data
- for unsupervised learning:
 - density estimation: larger (**structure match**) when $\mathbf{g}(\mathbf{x}) \approx \mathbf{x}$
 - outlier detection: those \mathbf{x} where $\mathbf{g}(\mathbf{x}) \not\approx \mathbf{x}$—learning **‘typical’ representation** of data

autoencoder:

representation-learning through
approximating identity function

Basic Autoencoder

basic **autoencoder**:

$d \rightarrow \tilde{d} \rightarrow d$ **NNet** with error function $\sum_{i=1}^d (g_i(\mathbf{x}) - x_i)^2$


- backprop **easily** applies; **shallow** and **easy** to train
- usually $\tilde{d} < d$: **compressed** representation
- data: $\{(\mathbf{x}_1, \mathbf{y}_1 = \mathbf{x}_1), (\mathbf{x}_2, \mathbf{y}_2 = \mathbf{x}_2), \dots, (\mathbf{x}_N, \mathbf{y}_N = \mathbf{x}_N)\}$
—often categorized as **unsupervised learning technique**
- sometimes constrain $w_{ij}^{(1)} = w_{ji}^{(2)}$ as **regularization**
—more **sophisticated** in calculating gradient

basic **autoencoder** in basic deep learning:
 $\{w_{ij}^{(1)}\}$ taken as **shallowly pre-trained weights**

Pre-Training with Autoencoders

Deep Learning with Autoencoders

- 1 for $\ell = 1, \dots, L$, **pre-train** $\{w_{ij}^{(\ell)}\}$ assuming $w_*^{(1)}, \dots, w_*^{(\ell-1)}$ fixed



by **training basic autoencoder** on $\{\mathbf{x}_n^{(\ell-1)}\}$ with $\tilde{d} = d^{(\ell)}$

- 2 **train with backprop** on **pre-trained** NNet to **fine-tune** all $\{w_{ij}^{(\ell)}\}$

many successful **pre-training** techniques take
'fancier' autoencoders with different
architectures and **regularization schemes**

Fun Time

Suppose training a d - \tilde{d} - d autoencoder with backprop takes approximately $c \cdot d \cdot \tilde{d}$ seconds. Then, what is the total number of seconds needed for pre-training a d - $d^{(1)}$ - $d^{(2)}$ - $d^{(3)}$ -1 deep NNet?

- ① $c (d + d^{(1)} + d^{(2)} + d^{(3)} + 1)$
- ② $c (d \cdot d^{(1)} \cdot d^{(2)} \cdot d^{(3)} \cdot 1)$
- ③ $c (dd^{(1)} + d^{(1)}d^{(2)} + d^{(2)}d^{(3)} + d^{(3)})$
- ④ $c (dd^{(1)} \cdot d^{(1)}d^{(2)} \cdot d^{(2)}d^{(3)} \cdot d^{(3)})$

Fun Time


Suppose training a d - \tilde{d} - d autoencoder with backprop takes approximately $c \cdot d \cdot \tilde{d}$ seconds. Then, what is the total number of seconds needed for pre-training a d - $d^{(1)}$ - $d^{(2)}$ - $d^{(3)}$ -1 deep NNet?

- ① $c (d + d^{(1)} + d^{(2)} + d^{(3)} + 1)$
- ② $c (d \cdot d^{(1)} \cdot d^{(2)} \cdot d^{(3)} \cdot 1)$
- ③ $c (dd^{(1)} + d^{(1)}d^{(2)} + d^{(2)}d^{(3)} + d^{(3)})$
- ④ $c (dd^{(1)} \cdot d^{(1)}d^{(2)} \cdot d^{(2)}d^{(3)} \cdot d^{(3)})$

Reference Answer: ③

Each $c \cdot d^{(\ell-1)} \cdot d^{(\ell)}$ represents the time for pre-training with one autoencoder to determine one layer of the weights.

Regularization in Deep Learning



watch out for overfitting, remember? :-)


high **model complexity**: **regularization** needed

- structural decisions/**constraints**
- weight decay or weight elimination **regularizers**
- **early stopping**

next: another **regularization** technique

Reasons of Overfitting Revisited

stochastic noise



reasons of serious overfitting:

data size N ↓	overfit ↑
noise ↑	overfit ↑
excessive power ↑	overfit ↑

how to deal with **noise**?

Dealing with Noise

- direct possibility: **data cleaning/pruning, remember? :-)**
- a **wild** possibility: **adding noise** to data?

- idea: **robust** autoencoder should not only let $\mathbf{g}(\mathbf{x}) \approx \mathbf{x}$ but also allow $\mathbf{g}(\tilde{\mathbf{x}}) \approx \mathbf{x}$ even when $\tilde{\mathbf{x}}$ slightly different from \mathbf{x}
- **denoising** autoencoder:

run basic autoencoder with data
 $\{(\tilde{\mathbf{x}}_1, \mathbf{y}_1 = \mathbf{x}_1), (\tilde{\mathbf{x}}_2, \mathbf{y}_2 = \mathbf{x}_2), \dots, (\tilde{\mathbf{x}}_N, \mathbf{y}_N = \mathbf{x}_N)\}$,
where $\tilde{\mathbf{x}}_n = \mathbf{x}_n + \text{artificial noise}$

—often used **instead of basic autoencoder** in deep learning

- useful for data/image processing: $\mathbf{g}(\tilde{\mathbf{x}})$ a **denoised** version of $\tilde{\mathbf{x}}$
- effect: ‘constrain/**regularize**’ \mathbf{g} towards **noise-tolerant** denoising

artificial noise/hint as **regularization**!
—practically also useful for other NNet/models

Fun Time

Which of the following cannot be viewed as a regularization technique?

- ① hint the model with artificially-generated noisy data
- ② stop gradient descent early
- ③ add a weight elimination regularizer
- ④ all the above are regularization techniques

Fun Time

Which of the following cannot be viewed as a regularization technique?

- ① hint the model with artificially-generated noisy data
- ② stop gradient descent early
- ③ add a weight elimination regularizer
- ④ all the above are regularization techniques

Reference Answer: ④

① is our new friend for regularization, while ② and ③ are old friends.

Linear Autoencoder Hypothesis

nonlinear autoencoder

sophisticated

linear autoencoder

simple

linear: more efficient? less overfitting? **linear first, remember? :-)**

linear hypothesis for k -th component
$$h_k(\mathbf{x}) = \sum_{j=0}^{\tilde{d}} w_{kj} \left(\sum_{i=1}^d w_{ij} x_i \right)$$

consider three special conditions:

- **exclude** x_0 : range of i **same** as range of k
- constrain $w_{ij}^{(1)} = w_{ji}^{(2)} = w_{ij}$: **regularization**
—denote $\mathbf{W} = [w_{ij}]$ of size $d \times \tilde{d}$
- assume $\tilde{d} < d$: ensure **non-trivial** solution

linear autoencoder hypothesis:

$$\mathbf{h}(\mathbf{x}) = \mathbf{W}\mathbf{W}^T \mathbf{x}$$

Linear Autoencoder Error Function

$$E_{\text{in}}(\mathbf{h}) = E_{\text{in}}(\mathbf{W}) = \frac{1}{N} \sum_{n=1}^N \left\| \mathbf{x}_n - \mathbf{W}\mathbf{W}^T \mathbf{x}_n \right\|^2 \text{ with } d \times \tilde{d} \text{ matrix } \mathbf{W}$$

—analytic solution to minimize E_{in} ? but **4-th order polynomial of w_{ij}**

let's familiarize the problem with linear algebra (**be brave! :-)**)

- eigen-decompose $\mathbf{W}\mathbf{W}^T = \mathbf{V}\mathbf{\Gamma}\mathbf{V}^T$
 - $d \times d$ matrix \mathbf{V} **orthogonal**: $\mathbf{V}\mathbf{V}^T = \mathbf{V}^T\mathbf{V} = \mathbf{I}_d$
 - $d \times d$ matrix $\mathbf{\Gamma}$ **diagonal** with $\leq \tilde{d}$ non-zero
- $\mathbf{W}\mathbf{W}^T \mathbf{x}_n = \mathbf{V}\mathbf{\Gamma}\mathbf{V}^T \mathbf{x}_n$
 - $\mathbf{V}^T(\mathbf{x}_n)$: change of **orthonormal basis** (**rotate** or reflect)
 - $\mathbf{\Gamma}(\cdots)$: set $\geq d - \tilde{d}$ components to 0, and **scale** others
 - $\mathbf{V}(\cdots)$: reconstruct by coefficients and **basis** (**back-rotate**)
- $\mathbf{x}_n = \mathbf{V}\mathbf{V}^T \mathbf{x}_n$: **rotate** and **back-rotate** cancel out

next: minimize E_{in} **by optimizing $\mathbf{\Gamma}$ and \mathbf{V}**

The Optimal Γ

$$\min_{\mathbf{V}} \min_{\Gamma} \frac{1}{N} \sum_{n=1}^N \left\| \underbrace{\mathbf{V} \mathbf{I} \mathbf{V}^T \mathbf{x}_n}_{\mathbf{x}_n} - \underbrace{\mathbf{V} \Gamma \mathbf{V}^T \mathbf{x}_n}_{\mathbf{W} \mathbf{W}^T \mathbf{x}_n} \right\|^2$$

- **back-rotate** not affecting length: ✗
- $\min_{\Gamma} \sum \|(I - \Gamma)(\text{some vector})\|^2$: **want many 0** within $(I - \Gamma)$
- optimal diagonal Γ with rank $\leq \tilde{d}$:

$$\left\{ \begin{array}{l} \tilde{d} \text{ diagonal components } 1 \\ \text{other components } 0 \end{array} \right\} \Rightarrow \text{without loss of gen. } \begin{bmatrix} \mathbf{I}_{\tilde{d}} & 0 \\ 0 & 0 \end{bmatrix}$$

$$\text{next: } \min_{\mathbf{V}} \sum_{n=1}^N \left\| \underbrace{\begin{bmatrix} 0 & 0 \\ 0 & \mathbf{I}_{d-\tilde{d}} \end{bmatrix}}_{\mathbf{I} - \text{optimal } \Gamma} \mathbf{V}^T \mathbf{x}_n \right\|^2$$

The Optimal \mathbf{V}

$$\min_{\mathbf{V}} \sum_{n=1}^N \left\| \begin{bmatrix} 0 & 0 \\ 0 & \mathbf{I}_{d-\tilde{d}} \end{bmatrix} \mathbf{V}^T \mathbf{x}_n \right\|^2 \equiv \max_{\mathbf{V}} \sum_{n=1}^N \left\| \begin{bmatrix} \mathbf{I}_{\tilde{d}} & 0 \\ 0 & 0 \end{bmatrix} \mathbf{V}^T \mathbf{x}_n \right\|^2$$

- $\tilde{d} = 1$: only first row \mathbf{v}^T of \mathbf{V}^T matters

$$\max_{\mathbf{v}} \sum_{n=1}^N \mathbf{v}^T \mathbf{x}_n \mathbf{x}_n^T \mathbf{v} \text{ subject to } \mathbf{v}^T \mathbf{v} = 1$$

- optimal \mathbf{v} satisfies $\sum_{n=1}^N \mathbf{x}_n \mathbf{x}_n^T \mathbf{v} = \lambda \mathbf{v}$
—using Lagrange multiplier λ , remember? :-)
- optimal \mathbf{v} : ‘topmost’ eigenvector of $\mathbf{X}^T \mathbf{X}$
- general \tilde{d} : $\{\mathbf{v}_j\}_{j=1}^{\tilde{d}}$ ‘topmost’ eigenvectors of $\mathbf{X}^T \mathbf{X}$
—optimal $\{\mathbf{w}_j\} = \{\mathbf{v}_j \text{ with } [\gamma_j = 1]\} = \text{top eigenvectors}$

linear autoencoder: projecting to orthogonal
patterns \mathbf{w}_j that ‘matches’ $\{\mathbf{x}_n\}$ most

Principal Component Analysis

Linear Autoencoder or **PCA**

- 1 let $\bar{\mathbf{x}} = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$, and let $\mathbf{x}_n \leftarrow \mathbf{x}_n - \bar{\mathbf{x}}$
- 2 calculate \tilde{d} top eigenvectors $\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_{\tilde{d}}$ of $\mathbf{X}^T \mathbf{X}$
- 3 return feature transform $\Phi(\mathbf{x}) = \mathbf{W}(\mathbf{x} - \bar{\mathbf{x}})$

- linear autoencoder:
maximize $\sum (\text{magnitude after projection})^2$
- **principal component analysis (PCA)** from statistics:
maximize $\sum (\text{variance after projection})$
- both useful for **linear dimension reduction**
though **PCA more popular**

linear dimension reduction:
useful for **data processing**

Fun Time

When solving the optimization problem

$$\max_{\mathbf{v}} \sum_{n=1}^N \mathbf{v}^T \mathbf{x}_n \mathbf{x}_n^T \mathbf{v} \text{ subject to } \mathbf{v}^T \mathbf{v} = 1,$$

we know that the optimal \mathbf{v} is the ‘topmost’ eigenvector that corresponds to the ‘topmost’ eigenvalue λ of $\mathbf{X}^T \mathbf{X}$. Then, what is the optimal objective value of the optimization problem?

- 1 λ^1
- 2 λ^2
- 3 λ^3
- 4 λ^4

Fun Time

When solving the optimization problem

$$\max_{\mathbf{v}} \sum_{n=1}^N \mathbf{v}^T \mathbf{x}_n \mathbf{x}_n^T \mathbf{v} \text{ subject to } \mathbf{v}^T \mathbf{v} = 1,$$

we know that the optimal \mathbf{v} is the ‘topmost’ eigenvector that corresponds to the ‘topmost’ eigenvalue λ of $\mathbf{X}^T \mathbf{X}$. Then, what is the optimal objective value of the optimization problem?

- ① λ^1
- ② λ^2
- ③ λ^3
- ④ λ^4

Reference Answer: ①

The objective value of the optimization problem is simply $\mathbf{v}^T \mathbf{X}^T \mathbf{X} \mathbf{v}$, which is $\lambda \mathbf{v}^T \mathbf{v}$ and you know what $\mathbf{v}^T \mathbf{v}$ must be.

Summary

- 1 Embedding Numerous Features: Kernel Models
- 2 Combining Predictive Features: Aggregation Models
- 3 Distilling Implicit Features: Extraction Models

Lecture 13: Deep Learning

- Deep Neural Network
hierarchical feature extraction
 - Autoencoder
unsupervised NNet learning of representation
 - Denoising Autoencoder
using noise as hints for regularization
 - Principal Component Analysis
linear autoencoder variant
- **next: extracting 'prototype' instead of pattern**